

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report, accuracy_score, precision_score, recall_score, f1_score
```

```
data = pd.read_csv('Social_Network_Ads.csv')
data.head(5)
```

```

  Age  EstimatedSalary  Purchased
0   19             19000           0
1   35             20000           0
2   26             43000           0
3   27             57000           0
4   19             76000           0
```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Age              400 non-null   int64  
1   EstimatedSalary  400 non-null   int64  
2   Purchased        400 non-null   int64  
dtypes: int64(3)
memory usage: 9.5 KB
```

```
data.describe()
```

```

  Age  EstimatedSalary  Purchased
count  400.000000      400.000000  400.000000
mean    37.655000      69742.500000  0.357500
std     10.482877      34096.960282  0.479864
min     18.000000      15000.000000  0.000000
25%     29.750000      43000.000000  0.000000
50%     37.000000      70000.000000  0.000000
75%     46.000000      88000.000000  1.000000
max     60.000000     150000.000000  1.000000
```

```
data.isnull().sum()
```

```

Age              0
EstimatedSalary  0
Purchased        0
dtype: int64
```

```
data.shape
```

```
(400, 3)
```

```

x = data.iloc[:, :2] # First two columns: Age and EstimatedSalary
y = data.iloc[:, 2]  # Third column: Purchased
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
```

```

scale = StandardScaler()
x_train = scale.fit_transform(x_train)
```

```
x_test = scale.transform(x_test)
```

```
lr = LogisticRegression(random_state = 0,solver = 'lbfgs')
lr.fit(x_train,y_train)
pred = lr.predict(x_test)
```

```
print(x_test[:10])
print('-'*15)
print(pred[:10])
```

```
[[ 0.812419 -1.39920777]
 [ 2.0889839  0.52871943]
 [-0.95513241 -0.75656537]
 [ 1.0088136  0.76240757]
 [-0.85693511 -1.22394166]
 [-0.75873781 -0.23076704]
 [ 0.9106163  1.08372877]
 [-0.85693511  0.38266434]
 [ 0.2232352  0.14897619]
 [ 0.4196298 -0.14313399]]
-----
[0 1 0 1 0 0 1 0 0 0]
```

```
print('Expected Output:',pred[:10])
print('-'*15)
print('Predicted Output:\n',y_test[:10])
```


```
Expected Output: [0 1 0 1 0 0 1 0 0 0]
-----
Predicted Output:
209    0
280    1
33     0
210    1
93     0
84     0
329    1
94     0
266    0
126    0
Name: Purchased, dtype: int64
```

```
matrix = confusion_matrix(y_test,pred,labels = lr.classes_)
print(matrix)
```

```
tp, fn, fp, tn = confusion_matrix(y_test,pred,labels=[1,0]).reshape(-1)
```

```
[[61  2]
 [12 25]]
```

```
conf_matrix = ConfusionMatrixDisplay(confusion_matrix=matrix,display_labels=lr.classes_)
conf_matrix.plot(cmap=plt.cm.Blues)
plt.show()
```



```
print(classification_report(y_test,pred))
```

```

precision    recall  f1-score   support

0           0.84        0.97        0.90         63
1           0.93        0.68        0.78         37

 accuracy          0.86         100
 macro avg          0.88        0.82        0.84         100
 weighted avg          0.87        0.86        0.85         100

```

```

print('\nAccuracy: {:.2f}'.format(accuracy_score(y_test,pred)))
print('Error Rate: ',(fp+fn)/(tp+tn+fn+fp))
print('Sensitivity (Recall or True positive rate) :',tp/(tp+fn))
print('Specificity (True negative rate) :',tn/(fp+tn))
print('Precision (Positive predictive value) :',tp/(tp+fp))
print('False Positive Rate :',fp/(tn+fp))

```

```

Accuracy: 0.86
Error Rate: 0.14
Sensitivity (Recall or True positive rate) : 0.6756756756756757
Specificity (True negative rate) : 0.9682539682539683
Precision (Positive predictive value) : 0.9259259259259259
False Positive Rate : 0.031746031746031744

```